

MPI和OpenMP简介

王政力


中科院理论物理研究所

wangzhengli@itp.ac.cn


2020年9月23日

问题描述

求和



扫点



Outline


- MPI
- OpenMP

什么是MPI

- **M**essage **P**assing **I**nterface (MPI) 是一种消息传递模型，由MPI Forum开发，官方支持C/C++，Fortran，社区支持Python，Ruby，Java，Julia，D等等
- MPI是一种库描述，并不是一种语言，有上百个函数调用接口，可以在C/C++，Fortran中直接对这些函数进行调用
- MPI是一种标准或者规范，并不是某种特指的具体实现

MPI的发展历程

- 1994年5月： MPI-1.0
 - 1995年6月： MPI-1.1
 - 1997年7月： MPI-1.2
 - 2008年5月： MPI-1.3
- 1998年： MPI-2.0
 - 2008年9月： MPI-2.1
 - 2009年9月： MPI-2.2
- 2012年9月： MPI-3.0
 - 2015年6月： MPI-3.1
- 目前： MPI-4.0正在开发中



Portability, efficiency, functionality

常见的MPI实现

- 主要有两大类库：Open MPI和MPICH（MPICH/MPICH2/MPICH3，OSU MVAPICH/MVAPICH2，Intel MPI，IBM MPI，Cray MPI）
- Intel MPI，Open MPI和MPICH3：支持InfiniBand和以太网
- MPICH/MPICH2：支持以太网，不支持InfiniBand
- MVAPICH/MVAPICH2：支持InfiniBand和以太网
- 区别：Open MPI的设计针对普遍通用的情况；MPICH跟进最新的MPI标准，并为其他有特殊需求的MPI实现提供了高质量的参考
- <https://stackoverflow.com/questions/2427399/mpich-vs-openmpi>

MPI基本概念

- 消息（变量）： `address, length`
- 源/目的地： `source/destination`，指明消息发往哪个进程或者从哪个进程接收消息
- 标签： `tag`，标记消息，选择性地发送或者接受指定的消息

`Send(address, length, destination, tag)`

`Receive(address, length, source, tag, actlen)`

MPI基本概念

- **MPI_Send**(address, count, datatype, destination, tag, comm)
- **MPI_Recv**(address, maxcount, datatype, source, tag, comm, status)
 - (address, count, datatype): (A, 300, **MPI_REAL**)代表Fortran中长度为300的实数数组
 - Comm: 通信域，MPI中最大的通信域为**MPI_COMM_WORLD**，进入并行环境时自动创建
 - Status: 接收到的消息的实际信息，包括消息的长度，来源和标签，接收进程可以接受特定来源特定标签的消息，也可以接受任意来源任意标签的消息，**MPI_ANY_SOURCE, MPI_ANY_TAG**

最基本的MPI函数

- **MPI_Init:** 初始化MPI
- **MPI_Comm_size:** 取得总进程数
- **MPI_Comm_rank:** 取得每个进程的进程号，从0开始
- **MPI_Send:** 发送消息
- **MPI_Recv:** 接收消息
- **MPI_Finalize:** 结束MPI

虽然MPI有上百个调用接口，但是其中最基本的六个函数就能编写一个完整的MPI程序，去解决很多问题

第一个MPI程序

各大编译器之间的mod文件并不兼容，
甚至同一编译器的不同版本也可能不兼容，
碰到编译错误可考虑更换编译器或者使用mpif.h
变量名千万不要取成mpi!!!

Fortran

```
1 program main
2   !use mpi
3   implicit none
4   include 'mpif.h'
5   integer :: ierr,myid,np
6   call mpi_init(ierr)
7   call mpi_comm_rank(mpi_comm_world,myid,ierr)
8   call mpi_comm_size(mpi_comm_world,np,ierr)
9   write(*,*) myid,np
10  call mpi_finalize(ierr)
11 end program main
```

C

```
1 #include "mpi.h"
2 #include <stdio.h>
3 int main(int argc, char *argv[]){
4   int myid,np;
5   MPI_Init(&argc,&argv);
6   MPI_Comm_rank(MPI_COMM_WORLD,&myid);
7   MPI_Comm_size(MPI_COMM_WORLD,&np);
8   printf("%d %d\n",myid,np);
9   MPI_Finalize();
10  return 0;
11 }
```


1	4
2	4
3	4
0	4

0	4
2	4
1	4
3	4

```

1 #include "mpi.h"
2 #include <stdio.h>
3 int main(int argc, char *argv[]){
4     int myid,np;
5     MPI_Init(&argc,&argv);
6     MPI_Comm_rank(MPI_COMM_WORLD,&myid);
7     MPI_Comm_size(MPI_COMM_WORLD,&np);
8     printf("%d %d\n",myid,np);
9     MPI_Finalize();
10    return 0;
11 }

```



SPMD: Single Program, Multiple Data

各个进程之间没有影响，每个语句独立地在各个进程中执行，即使是输出语句，所以如果输出很多，输出结果可能很乱

```


1 #include "mpi.h"
2 #include <stdio.h>
3 1 #include "mpi.h"
4 2 #include <stdio.h>
5 3 1 #include "mpi.h"
6 4 2 #include <stdio.h>
7 5 3 1 #include "mpi.h"
8 6 4 2 #include <stdio.h>
9 7 5 3 int main(int argc, char *argv[]){
10 8 6 4 int myid,np;
11 9 7 5 MPI_Init(&argc,&argv);
10 8 6 MPI_Comm_rank(MPI_COMM_WORLD,&myid);
11 9 7 MPI_Comm_size(MPI_COMM_WORLD,&np);
10 8 printf("%d %d\n",myid,np);
11 9 MPI_Finalize();
10 return 0;
11 }

```

```

0 4
2 4
1 4
3 4

```



MPI程序的编译

- Open MPI（不分编译器）

- C: mpicc
- C++: mpic++, mpicxx, mpiCC
- Fortran: mpifort(mpif77, mpif90)

- Intel MPI（分编译器）

- GNU编译器: mpicc, mpigcc, mpicxx, mpigxx, mpifc(mpif77, mpif90)
- Intel编译器: mpiicc, mpiicpc, mpiifort

- 这些命令只是一层封装，实际会调用系统编译器，编译时自动添加MPI头文件搜索路径，链接时自动链接MPI库文件，方便使用

```
mpicc -show
gcc -I/public/software//mpi/intelmpi/2017.4.239/intel64
intelmpi/2017.4.239/intel64/lib -Xlinker --enable-new-dt
r -rpath -Xlinker /public/software//mpi/intelmpi/2017.4
-rpath -Xlinker /opt/intel/mpi-rt/2017.0.0/intel64/lib
```

MPI程序的运行

- 常见的MPI实现都提供以下命令
 - `mpirun -np 4 ./a.out`
- MPI Forum建议使用以下命令
 - `mpiexec -n 4 ./a.out`

Fortran & C bindings

Fortran	C
MPI_INIT (ierror) integer ierror	int MPI_Init (int *argc, ***argv)
MPI_COMM_SIZE (comm, size, ierror) integer comm, size, ierror	int MPI_Comm_size (MPI_Comm comm, int *size)
MPI_COMM_RANK (comm, rank, ierror) integer comm, rank, ierror	int MPI_Comm_rank (MPI_Comm comm, int *rank)
MPI_FINALIZE (ierror) integer comm	int MPI_Finalize ()

第二个MPI程序：求和


$$f(x) = \frac{4}{1+x^2}$$

$$\Delta x = \frac{1}{n}$$

$$x_i = \frac{i - 0.5}{n}$$

$$\pi \approx \sum_{i=1}^n f(x_i) \Delta x$$

$$\begin{aligned} \int_0^1 \frac{1}{1+x^2} dx &= \arctan(x) \Big|_0^1 \\ &= \arctan(1) - \arctan(0) \\ &= \arctan(1) \\ &= \frac{\pi}{4} \end{aligned}$$



```

1 program main
2   use mpi
3   implicit none
4   integer, parameter :: dbl = selected_real_kind(15)
5   real(dbl), parameter :: pi = acos(-1.0_dbl)
6   integer :: ierr, myid, np, i, n
7   real(dbl) :: dx, xi, mypi, tpi
8   call mpi_init(ierr)
9   call mpi_comm_rank(mpi_comm_world, myid, ierr)
10  call mpi_comm_size(mpi_comm_world, np, ierr)
11  do
12    if ( myid == 0 ) then
13      write(*,*) "Enter the number of intervals:(0 quits)"
14      read(*,*) n
15    end if
16    call mpi_bcast(n, 1, mpi_integer, 0, mpi_comm_world, ierr)
17    if ( n <= 0 ) exit
18    dx = 1.0_dbl/n
19    mypi = 0.0_dbl
20    do i = myid+1, n, np
21      xi = (i-0.5_dbl)/n
22      mypi = mypi + 4.0_dbl/(1.0_dbl + xi**2)*dx
23    end do
24    call mpi_reduce(mypi, tpi, 1, mpi_double_precision, &
25      mpi_sum, 0, mpi_comm_world, ierr)
26    if ( myid == 0 ) then
27      write(*,*) pi
28      write(*,*) tpi
29      write(*,*) abs(pi-tpi)/pi*100, "%"
30    end if
31  end do
32  call mpi_finalize(ierr)
33 end program main

```

```

1 #include "mpi.h"
2 #include <stdio.h>
3 #include <math.h>
4 void main(int argc, char *argv[]) {
5   const double pi = acos(-1.0);
6   double dx, xi, mypi, tpi;
7   int myid, np, n, i;
8   MPI_Init( &argc, &argv);
9   MPI_Comm_rank( MPI_COMM_WORLD, &myid);
10  MPI_Comm_size( MPI_COMM_WORLD, &np);
11  while (1) {
12    if (myid == 0) {
13      printf("Enter the number of intervals:(0 quits)\n");
14      scanf("%d", &n);
15    }
16    MPI_Bcast( &n, 1, MPI_INT, 0, MPI_COMM_WORLD);
17    if (n<=0) break;
18    dx = 1.0/n;
19    mypi = 0.0;
20    for (i = myid+1; i <= n; i+=np) {
21      xi = (i-0.5)/n;
22      mypi += 4.0/(1.0+xi*xi)*dx;
23    }
24    MPI_Reduce( &mypi, &tpi, 1, MPI_DOUBLE, MPI_SUM, 0,
25      MPI_COMM_WORLD);
26    if (myid==0) {
27      printf("%.16f\n", pi);
28      printf("%.16f\n", tpi);
29      printf("%.16e%%\n", fabs(pi-tpi)/pi*100);
30    }
31  }
32  MPI_Finalize();
33 }

```


Before MPI_Bcast

Process 1

10

Process 2

Process 3

Process 4

After MPI_Bcast

Process 1

10

Process 2

10

Process 3

10

Process 4

10

Before MPI_Reduce

Process 1

1

Process 2

2

Process 3

3

Process 4

4

After MPI_Reduce

Process 1

10

Process 2

Process 3

Process 4

```
Enter the number of intervals:(0 quits)
9
3.14159265358979
3.14262145655761
3.274781555922022E-002 %
Enter the number of intervals:(0 quits)
99
3.14159265358979
3.14160115612355
2.706440551352709E-004 %
Enter the number of intervals:(0 quits)
999
3.14159265358979
3.14159273709004
2.657895530102677E-006 %
Enter the number of intervals:(0 quits)
9999
3.14159265358979
3.14159265442329
2.653108456047908E-008 %
Enter the number of intervals:(0 quits)
99999
3.14159265358979
3.14159265359813
2.653006678298102E-010 %
Enter the number of intervals:(0 quits)
0
```

```
Enter the number of intervals:(0 quits)
9
3.1415926535897931
3.1426214565576127
3.2747815559220218e-02%
Enter the number of intervals:(0 quits)
99
3.1415926535897931
3.1416011561235466
2.7064405513527092e-04%
Enter the number of intervals:(0 quits)
999
3.1415926535897931
3.1415927370900434
2.6578955159668783e-06%
Enter the number of intervals:(0 quits)
9999
3.1415926535897931
3.1415926544232935
2.6531141103673421e-08%
Enter the number of intervals:(0 quits)
99999
3.1415926535897931
3.1415926535981287
2.6532893942697868e-10%
Enter the number of intervals:(0 quits)
0
```

Fortran & C bindings

Fortran	MPI_BCAST (buffer, count, datatype, root, comm, ierror) <type> buffer(*) integer count, datatype, root, comm, ierror
	MPI_REDUCE (sendbuf, recvbuf, count, datatype, op, root, comm, ierror) <type> sendbuf(*), recvbuf(*) integer count, datatype, op, root, comm, ierror
C	int MPI_Bcast (void *buf, int count, MPI_Datatype datatype, int root, MPI_Comm comm)
	int MPI_Reduce (const void *sendbuf, void *recvbuf, int count, MPI_Datatype, MPI_Op op, int root, MPI_Comm comm)


MPI归约操作

MPI_MAX	最大值
MPI_MIN	最小值
MPI_MAXLOC	最大值且相应位置
MPI_MINLOC	最小值且相应位置
MPI_SUM	求和
MPI_PROD	乘积

MPI基本数据类型

MPI Fortran		MPI C	
MPI_CHARACTER	character(1)	MPI_CHAR	signed char
MPI_INTEGER	integer	MPI_INT	signed int
MPI_REAL	real	MPI_FLOAT	float
MPI_DOUBLE_PRECISION	double precision	MPI_DOUBLE	double
MPI_COMPLEX	complex	MPI_LONG	signed long int
MPI_DOUBLE_COMPLEX	double complex	MPI_LONG_DOUBLE	long double
MPI_LOGICAL	logical	MPI_UNSIGNED	unsigned int

第三个MPI程序： 矩阵矢量乘积



Self-scheduling or manager-worker algorithm

- 主要思想是一个进程负责分配任务，其余进程负责计算
- 这个机制非常适合其余进程（**worker**）不需要相互通信或者每个**worker**需要执行的任务量不确定的情况，所以scheduling可以理解为负载均衡
- 几个例子：
 - 矩阵乘积，计算每个矩阵元的计算量一样
 - 扫描T矩阵，有时候阈值上下计算量并不一样
 - 用试除法求N以内的素数个数，判断一个数是否为素数所需要的计算量并不相同

Common part

$$c = a \times b$$

```
1 program main
2   use mpi
3   implicit none
4   integer, parameter :: dbl = selected_real_kind(15)
5   integer, parameter :: n = 1000
6   real(dbl) :: a(n,n), b(n), c(n), buffer(n), ans
7   integer :: ierr, myid, np, manager, stat(mpi_status_size)
8   integer :: i, numsent, source, tag
9   call mpi_init(ierr)
10  call mpi_comm_rank(mpi_comm_world, myid, ierr)
11  call mpi_comm_size(mpi_comm_world, np, ierr)
12  manager = 0
13  if(myid == manager) then
14  else
15  end if
16  call mpi_finalize(ierr)
17 end program main
```

Manager part

- 先把b广播到各个进程
- 再把a的行发送到每个worker
- 进入循环，每接收一个c矩阵元就再发送a的一行到worker，直到填满c为止，这时候发送一个终止信号到各个worker

```
13  if(myid == manager) then
14      a = 1.0_dbl
15      b = 1.0_dbl
16      numsent = 0
17      call mpi_bcast(b,n,mpi_double_precision,&
18                    manager,mpi_comm_world,ierr)
19      do i = 1, np-1
20          buffer = a(i,:)
21          call mpi_send(buffer,n,mpi_double_precision,&
22                        i,i,mpi_comm_world,ierr)
23          numsent = numsent + 1
24      end do
25      do i = 1, n
26          call mpi_recv(ans,1,mpi_double_precision,&
27                        mpi_any_source,mpi_any_tag,mpi_comm_world,stat,ierr)
28          source = stat(mpi_source)
29          tag = stat(mpi_tag)
30          c(tag) = ans
31          if(numsent<n)then
32              buffer = a(numsent+1,:)
33              call mpi_send(buffer,n,mpi_double_precision,&
34                            source,numsent+1,mpi_comm_world,ierr)
35              numsent = numsent + 1
36          else
37              call mpi_send(mpi_bottom,0,mpi_double_precision,&
38                            source,0,mpi_comm_world,ierr)
39          endif
40      end do
41      write(*,*) maxval(abs(c-n))
42  else
```


Worker part

- 先得到b
- 进入循环，包括接收a的行，计算矩阵元，把结果发送给manager，直到收到终止信号结束循环

```
42  else
43      call mpi_bcast(b,n,mpi_double_precision,&
44                    manager,mpi_comm_world,ierr)
45  do
46      call mpi_recv(buffer,n,mpi_double_precision,&
47                    manager,mpi_any_tag,mpi_comm_world,stat,ierr)
48      tag = stat(mpi_tag)
49      if(tag == 0) exit
50      ans = dot_product(buffer,b)
51      call mpi_send(ans,1,mpi_double_precision,&
52                    manager,tag,mpi_comm_world,ierr)
53  end do
54  end if
```

几点说明

- 默认情况下tag的取值范围0-32767
- Status是输出参数，包含接收到的消息的source，tag，length
- 在Fortran中，status(**MPI_STATUS_SIZE**)是整数数组，status(**MPI_SOURCE**)为消息来源，status(**MPI_TAG**)为消息标签
- 在C中，status是**MPI_Status**类型的结构体，status.**MPI_SOURCE**表示消息的来源，status.**MPI_TAG**表示消息的标签

Fortran	MPI_SEND (buf, count, datatype, dest, tag, comm, ierror) <type> buf(*) integer count, datatype, dest, tag, comm, ierror
	MPI_RECV (buf, count, datatype, source, tag, comm, status, ierror) <type> buf(*) integer count, datatype, source, tag, comm, status(MPI_STATUS_SIZE), ierror
C	int MPI_Send (const void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
	int MPI_Recv (void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)

Outline

- MPI
- OpenMP

什么是OpenMP

- OpenMP(Open Multi-Processing)是一种针对共享式内存的多线程并行变成标准（不能像MPI那样跨节点），支持C/C++, Fortran
- 由编译指令，库函数，环境变量三部分组成，通常由编译器提供支持，在编译时打开OpenMP编译选项即可
- 已有的代码不需要大幅度修改，只需加上几条OpenMP指令即可并行化，当编译器不支持OpenMP时，程序又可退化为串行程序
- 行业标准：
 - DEC, Intel, IBM, HP, Sun, SGI, Fujitsu, 美国能源部等等
 - 跨平台：Windows/Linux/macOS

历史回顾

- 过去每个厂商都有自己的指令和库标准，虽然有ANSI X3H5标准，但未被正式采用；二来有消息传递库PVM和MPI，能够代替内存共享机
- 但在1996-1997年，厂商恢复了对共享内存编程架构的兴趣，还有部分厂商认为消息传递模型太过复杂
- 进入多核cpu时代后，用操作系统API编写多线程程序不如OpenMP来得方便，比如扩展性问题，可移植性问题等等

Release History

Date	Version
Oct 1997	Fortran 1.0
Oct 1998	C/C++ 1.0
Nov 1999	Fortran 1.1
Nov 2000	Fortran 2.0
Mar 2002	C/C++ 2.0
May 2005	OpenMP 2.5
May 2008	OpenMP 3.0
Jul 2011	OpenMP 3.1
Jul 2013	OpenMP 4.0
Nov 2015	OpenMP 4.5
Nov 2018	OpenMP 5.0

GNU	GCC C/C++/Fortran	<p>Free and open source – Linux, Solaris, AIX, MacOSX, Windows, FreeBSD, NetBSD, OpenBSD, DragonFly BSD, HPUX, RTEMS</p> <ul style="list-style-type: none"> ▪ From GCC 4.2.0, OpenMP 2.5 is fully supported for C/C++/Fortran. ▪ From GCC 4.4.0, OpenMP 3.0 is fully supported for C/C++/Fortran. ▪ From GCC 4.7.0, OpenMP 3.1 is fully supported for C/C++/Fortran. ▪ In GCC 4.9.0, OpenMP 4.0 is supported for C and C++, but not Fortran. ▪ From GCC 4.9.1, OpenMP 4.0 is fully supported for C/C++/Fortran. ▪ From GCC 6.1, OpenMP 4.5 is fully supported for C and C++. ▪ From GCC 7.1, OpenMP 4.5 is partially supported for Fortran. ▪ From GCC 9.1, OpenMP 5.0 is partially supported for C and C++. <p>Compile with -fopenmp to enable OpenMP.</p> <p>Online documentation: https://gcc.gnu.org/onlinedocs/libgomp/ OpenMP support history: https://gcc.gnu.org/projects/gomp/</p>
IBM	XL C/C++/Fortran	<p>XL C/C++ for Linux V16.1.1 and XL Fortran for Linux V16.1.1 fully support OpenMP 4.5 features including the target constructs.</p> <p>Compile with -qsmp=omp to enable OpenMP directives and with -qoffload for offloading the target regions to GPUs.</p> <p>For more information, please visit IBM XL C/C++ for Linux and IBM XL Fortran for Linux.</p>
Intel	C/C++/Fortran	<p>Windows, Linux, and MacOSX.</p> <ul style="list-style-type: none"> ▪ OpenMP 3.1 C/C++/Fortran fully supported in version 12.0, 13.0, 14.0 compilers ▪ OpenMP 4.0 C/C++/Fortran supported in version 15.0 and 16.0 compilers ▪ OpenMP 4.5 C/C++/Fortran supported in version 17.0, 18.0, and 19.0 compilers ▪ OpenMP 4.5 and subset of OpenMP 5.0 C/C++/Fortran supported in 19.1 compilers under -qnextgen -fiopenmp. <p>Compile with -Qopenmp on Windows, or just -openmp or -qopenmp on Linux or Mac OSX</p> <p>More information</p>

常用编译器的OpenMP编译选项

编译器	GNU	Intel	Nvidia（原PGI）
OpenMP选项	-fopenmp	-qopenmp	-mp

2013年7月29日，PGI被Nvidia收购

2020年8月05日，PGI成为NVIDIA HPC SDK的一部分，可以从Nvidia免费下载

第一个OpenMP程序

```
1 program main
2   implicit none
3   write(*,*) 'Hi'
4   !$omp parallel
5   write(*,*) 'Hello'
6   !$omp end parallel
7   write(*,*) 'Bye'
8 end program main
```

```
gfortran hello.f90 && ./a.out
Hi
Hello
Bye
```

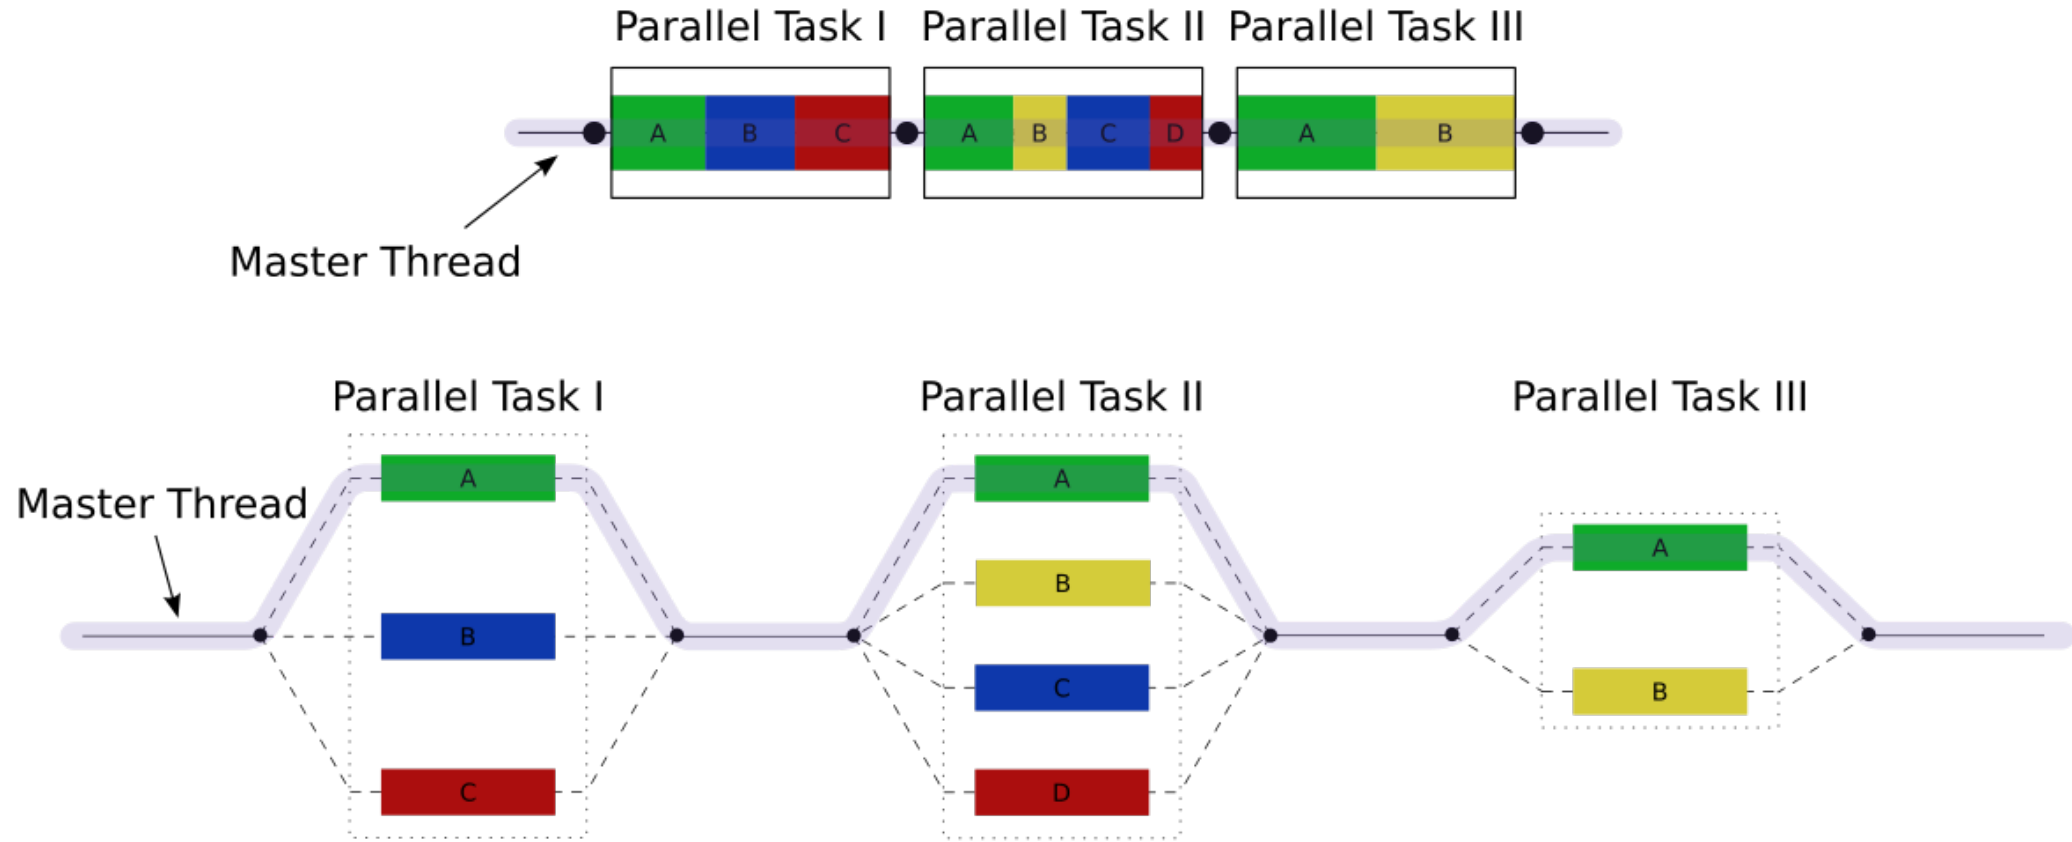
```
gfortran hello.f90 -fopenmp && ./a.out
Hi
Hello
Hello
Hello
Hello
Bye
```

```
1 #include <stdio.h>
2 void main(int argc, char *argv[]){
3   printf("Hi\n");
4   #pragma omp parallel
5   {
6     printf("Hello\n");
7   }
8   printf("Bye\n");
9 }
```

```
gcc hello.c && ./a.out
Hi
Hello
Bye
```

```
gcc hello.c -fopenmp && ./a.out
Hi
Hello
Hello
Hello
Hello
Bye
```

Fork-Join Model



OpenMP API Overview

- Compiler Directives

- Syntax: sentinel directive-name [clause, ...]
- Fortran: !\$OMP PARALLEL DEFAULT(SHARED) PRIVATE(BETA,PI)
- C/C++: #pragma omp parallel default(shared) private(beta,pi)

- Runtime Library Routines

- omp_get_thread_num, omp_get_num_threads
- Fortran: omp_lib.h(use omp_lib)
- C/C++: omp.h

- Environment Variables

- OMP_NUM_THREADS

Directives

Work-sharing constructs	Combined parallel work-sharing constructs	Synchronization constructs	Data environment constructs
do/for sections single	parallel do/parallel for parallel sections	master critical barrier atomic flush ordered	threadprivate

do/for

```
1 program main
2   implicit none
3   integer, parameter :: dbl = selected_real_kind(15)
4   integer, parameter :: n = 2000
5   real(dbl), dimension(n,n) :: a,b,c
6   integer :: i,j,k
7   a = 1.0_dbl
8   b = 1.0_dbl
9   c = 0.0_dbl
10  !$omp parallel
11  !$omp do
12  do i = 1, n
13      do j = 1, n
14          do k = 1, n
15              c(i,j) = c(i,j) + a(i,k)*b(k,j)
16          end do
17      end do
18  end do
19  !$omp end do
20  !$omp end parallel
21  write(*,*) maxval(abs(c-n))
22 end program main
```

```
1 #include <stdio.h>
2 void main(int argc, char *argv[]){
3     const int n = 2000;
4     double a[n][n],b[n][n],c[n][n];
5     for (int i=0; i<n; i++){
6         for (int j=0; j<n; j++){
7             a[i][j] = 1.0;
8             b[i][j] = 1.0;
9             c[i][j] = 0.0;
10        }
11    }
12    #pragma omp parallel
13    {
14        #pragma omp for
15        for (int i=0; i<n; i++){
16            for (int j=0; j<n; j++){
17                for (int k=0; k<n; k++){
18                    c[i][j] += a[i][k]*b[k][j];
19                }
20            }
21        }
22    }
23    printf("%f\n",c[0][0]-n);
24 }
```

```
gfortran -O2 mm.f90 && time ./a.out  
0.000000000000000000
```

```
real    0m12.518s  
user    0m12.492s  
sys     0m0.026s
```

```
gcc -O2 mm.c && time ./a.out  
0.000000
```

```
real    0m12.080s  
user    0m12.058s  
sys     0m0.022s
```

```
gfortran -O2 mm.f90 -fopenmp && time ./a.out  
0.000000000000000000
```

```
real    0m3.702s  
user    0m14.576s  
sys     0m0.031s
```

```
gcc -O2 mm.c -fopenmp && time ./a.out  
0.000000
```

```
real    0m3.688s  
user    0m14.303s  
sys     0m0.046s
```

sections/single/master

```
1 program main
2   implicit none
3   !$omp parallel sections
4   !$omp section
5   write(*,*) 'Hello'
6   !$omp section
7   write(*,*) 'Hi'
8   !$omp section
9   write(*,*) 'Bye'
10  !$omp end parallel sections
11 end program main
```

```
gfortran -fopenmp tmp.f90 && ./a.out
Hello
Hi
Bye
```

```
1 program main
2   implicit none
3   !$omp parallel
4   write(*,*) 'Hello'
5   !$omp single
6   write(*,*) 'Hi'
7   !$omp end single
8   !$omp end parallel
9 end program main
```

```
gfortran -fopenmp tmp.f90 && ./a.out
Hello
Hi
Hello
Hello
Hello
```

```
1 program main
2   implicit none
3   !$omp parallel
4   write(*,*) 'Hello'
5   !$omp master
6   write(*,*) 'Hi'
7   !$omp end master
8   !$omp end parallel
9 end program main
```

```
gfortran -fopenmp tmp.f90 && ./a.out
Hello
Hi
Hello
Hello
Hello
```

critical/atomic

```
1 program main
2   implicit none
3   integer, parameter :: n = 100
4   integer :: i,s
5   s = 0
6   !$omp parallel do
7   do i=1,n
8     !$omp critical
9     s = s + i
10    !$omp end critical
11  end do
12  !$omp end parallel do
13  write(*,*) s
14 end program main
gfortran -fopenmp tmp.f90 && ./a.out
5050
```

```
1 program main
2   implicit none
3   integer, parameter :: n = 100
4   integer :: i,s
5   s = 0
6   !$omp parallel do
7   do i=1,n
8     s = s + i
9   end do
10  !$omp end parallel do
11  write(*,*) s
12 end program main
gfortran -fopenmp tmp.f90 && ./a.out
3150
```

```
1 program main
2   implicit none
3   integer, parameter :: n = 100
4   integer :: i,s
5   s = 0
6   !$omp parallel do
7   do i=1,n
8     !$omp atomic
9     s = s + i
10  end do
11  !$omp end parallel do
12  write(*,*) s
13 end program main
gfortran -fopenmp tmp.f90 && ./a.out
5050
```


Clauses

Data scope attribute clauses	Other clauses
PRIVATE(list) SHARED(list) DEFAULT(PRIVATE SHARED NONE) FIRSTPRIVATE(list) LASTPRIVATE(list) COPYIN(list) COPYPRIVATE(list) REDUCTION(operator:list)	IF(scalar logical expression) NUM_THREADS(scalar integer expression) NOWAIT SCHEDULE(type, chunk)

Reduction(operator:list)

Operation	Fortran	C/C++	Initialization
Addition	+	+	0
Multiplication	*	*	1
Subtraction	-	-	0
Logical AND	.and.	&&	.true. / 1
Logical OR	.or.		.false. / 0
AND bitwise	land	&	All bits on / ~0
OR bitwise	lor		0
Exclusive OR bitwise	leor	^	0
Equivalent	.eqv.		.true.
Not Equivalent	.neqv.		.false.
Maximum	Max	Max	Most negative
Minimum	Min	Min	Largest positive

```

1 program main
2   implicit none
3   integer, parameter :: dbl = selected_real_kind(15)
4   integer, parameter :: n = 100
5   real(dbl), parameter :: pi = acos(-1.0_dbl)
6   real(dbl) :: mypi,dx,xi
7   integer :: i
8   mypi = 0.0_dbl
9   dx = 1.0_dbl/dbl(n)
10  !$omp parallel private(xi)
11  !$omp do reduction(+:mypi)
12  do i=1,n
13      xi = (dbl(i)-0.5_dbl)/dbl(n)
14      mypi = mypi + 4.0_dbl/(1.0_dbl+xi**2)*dx
15  end do
16  !$omp end do
17  !$omp end parallel
18  write(*,*) pi
19  write(*,*) mypi
20  write(*,*) abs(pi-mypi)/pi*100, '%'
21 end program main

```

```

gfortran -fopenmp pi.f90 && ./a.out
3.1415926535897931
3.1416009869231249
2.6525823843875473E-004 %

```

```

1 #include <stdio.h>
2 #include <math.h>
3 void main(int argc, char *argv[]){
4   const int n = 100;
5   const double pi = acos(-1.0);
6   double mypi,dx,xi;
7   mypi = 0.0;
8   dx = 1.0/(double)n;
9   #pragma omp parallel private(xi)
10  {
11    #pragma omp for reduction(+:mypi)
12    for (int i=1; i<=n; i++){
13        xi = ((double)i-0.5)/(double)n;
14        mypi += 4.0/(1.0+xi*xi)*dx;
15    }
16  }
17  printf("%f\n",pi);
18  printf("%f\n",mypi);
19  printf("%f%%\n",fabs(pi-mypi)/pi*100);
20 }

```

```

gcc -fopenmp pi.c && ./a.out
3.141593
3.141601
0.000265%

```

私有&公有变量

- 默认情况下循环计数器和**threadprivate**变量是私有变量，其余都是公有变量，私有变量在并行开始处其值不确定，并行结束时其值也不确定
- 通过**private**，**share**，**default**来声明变量属性
- **firstprivate**在并行开始处对私有变量赋值，**lastprivate**在并行结束时更新私有变量的值，**copyprivate**可在单一线程执行**single**指令对中的指令后将私有变量的值传递给其他线程

Schedule(type, chunk)

- 并行执行循环语句时，每个线程执行相同数量的迭代，这并不总是最好的办法，因为每个线程的工作量可能会不相同
- **Static:** chunk不赋值时，默认每个线程执行相同数量的迭代，比如1-25, 26-50, 51-75, 76-100
- **Dynamic:** chunk不赋值时，默认chunk=1，每个线程执行一次迭代，比如1, 2, 3, 4，最先执行完的线程继续执行下一个迭代
- **Guided:** 与静态方法相比，动态方法提高了执行能力和效率，但处理和分配任务时产生了冗余，工作块越小，冗余越多；另一种动态方法是每个块的工作量按指数递减，后一个块的工作量是前一个的一半
- **Runtime:** 通过环境变量OMP_SCHEDULE来指定

static

```
gfortran -O2 prime.f90 -fopenmp && time ./a.out
5761455

real    0m19.608s
user    0m56.780s
sys     0m0.176s
```

dynamic

```
gfortran -O2 prime.f90 -fopenmp && time ./a.out
5761455

real    0m15.105s
user    0m59.481s
sys     0m0.162s
```

guided

```
gfortran -O2 prime.f90 -fopenmp && time ./a.out
5761455

real    0m14.326s
user    0m56.540s
sys     0m0.151s
```

```
1 function p(n)
2   implicit none
3   integer :: n,p,i
4   do i=3,nint(sqrt(dble(n))),2
5     if(mod(n,i)==0) then
6       p = 0
7       return
8     end if
9   end do
10  p = 1
11 end function p
12
13 program main
14   implicit none
15   integer, parameter :: n = 1e8
16   integer :: a(n),i
17   integer, external :: p
18   a = 0
19   !$omp parallel do schedule(dynamic)
20   do i=3,n,2
21     a(i) = p(i)
22   end do
23   !$omp end parallel do
24   write(*,*) sum(a)+1
25 end program main
```

num_threads & if

```
1 program main
2   implicit none
3   !$omp parallel num_threads(3)
4   !$omp sections
5   !$omp section
6   write(*,*) 'Hi'
7   !$omp section
8   write(*,*) 'Hello'
9   !$omp section
10  write(*,*) 'Bye'
11  !$omp end sections
12  !$omp end parallel
13 end program main
```

```
1 function p(n)
2   implicit none
3   integer :: n,p,i
4   do i=3,nint(sqrt(dble(n))),2
5     if(mod(n,i)==0) then
6       p = 0
7       return
8     end if
9   end do
10  p = 1
11 end function p
12
13 program main
14   implicit none
15   integer, parameter :: n = 1e8
16   integer :: a(n),i
17   integer, external :: p
18   a = 0
19   !$omp parallel if(n>1e5)
20   !$omp do schedule(dynamic)
21   do i=3,n,2
22     a(i) = p(i)
23   end do
24   !$omp end do
25   !$omp end parallel
26   write(*,*) sum(a)+1
27 end program main
```

库函数

Execution environment routines

OMP_set_num_threads

OMP_get_num_threads

OMP_get_max_threads

OMP_get_thread_num

OMP_get_num_procs

OMP_in_parallel

OMP_set_dynamic

OMP_get_dynamic

OMP_set_nested

OMP_get_nested

```
1 program main
2   !use omp_lib
3   implicit none
4   include 'omp_lib.h'
5   integer,save :: myid,np
6   !$omp threadprivate(myid,np)
7   !$omp parallel
8   myid = omp_get_thread_num()
9   np = omp_get_num_threads()
10  !$omp end parallel
11  !$omp parallel
12  write(*,*) myid,np
13  !$omp end parallel
14 end program main
```

```
gfortran -fopenmp tmp.f90 && ./a.out
```

```
2      4
0      4
1      4
3      4
```


环境变量

The environment variables

OMP_NUM_THREADS

OMP_SCHEDULE

OMP_DYNAMIC

OMP_NESTED

优先级:

NUM_THREADS >

OMP_SET_NUM_THREADS >

OMP_NUM_THREADS

总结

- MPI: 进程级, 分布式存储, 效率高, 可扩展性好, 程序复杂
- OpenMP: 线程级, 共享式存储, 可扩展性差 (单机, 不可跨节点), 程序简单, 要注意racing condition

References

- Using MPI Portable Parallel Programming with the Message-Passing Interface
- Using MPI-2 Advanced Features of the Message-passing Interface
- Parallel Programming in Fortran 95 using OpenMP

谢谢！